

Data

Below is the table with our findings. The units are in seconds for all tests. The data was obtained by logging the “real” value from the time command. The average value of 2 tests was recorded for every test case, and those values was used for the speedup and efficiency tables.

The tests were interleaved to leave at least 10 minutes between each test for the same case.

	Cores	Image Size			
		1920x1080	2560x1600	2880x2560	3840x2160
Radius 0	1	6.8645	12.9965	22.9335	25.8835
	2	3.7625	6.8825	11.903	13.317
	4	2.2145	3.82	6.3975	7.125
	8	1.3825	2.4145	3.63	4.017
	16	1.1455	1.444	2.2455	2.487
	32	0.933	1.001	1.305	1.8615
Radius 10	1	983.347	1751.091	2216.57	3944.869
	2	493.486	875.295	1108.487	1972.151
	4	247.444	440.055	556.795	990.075
	8	124.121	220.292	278.027	495.305
	16	66.16	110.541	139.626	248.061
	32	49.995	55.566	92.3345	146.4735
Radius 20	32	177.356	207.712	332.6395	531.9305
Radius 40	32	623.852	816.765	1212.836	1985.175

At the time of testing, mpirun was not working. The serial execution time was instead recorded by running the CUDA code with one process only.

Speedup					
	Cores	1920x1080	2560x1600	2880x2560	3840x2160
Radius 0	1	1	1	1	1
	2	1.824452	1.88833999	1.926699	1.943643
	4	3.099797	3.40222513	3.58476	3.632772
	8	4.96528	5.38268793	6.317769	6.44349
	16	5.99258	9.00034626	10.21309	10.40752
	32	7.357449	12.9835165	17.57356	13.90465
Radius 10	1	1	1	1	1
	2	1.992654	2.00057238	1.999636	2.000288
	4	3.974018	3.97925487	3.980945	3.984414
	8	7.922487	7.94895412	7.972499	7.964525
	16	14.86317	15.8410997	15.87505	15.90282
	32	19.66537	31.5159821	23.99976	26.92982

Efficiency					
	Cores	1920x1080	2560x1600	2880x2560	3840x2160
Radius 0	1	1	1	1	1
	2	0.912226	0.94417	0.96335	0.971822
	4	0.774949	0.85055628	0.89619	0.908193
	8	0.62066	0.67283599	0.789721	0.805436
	16	0.374536	0.56252164	0.638318	0.65047
	32	0.22992	0.40573489	0.549174	0.43452
Radius 10	1	1	1	1	1
	2	0.996327	1.00028619	0.999818	1.000144
	4	0.993505	0.99481372	0.995236	0.996104
	8	0.990311	0.99361926	0.996562	0.995566
	16	0.928948	0.99006873	0.992191	0.993926
	32	0.614543	0.98487444	0.749993	0.841557

Discussion

There was very little variation between the runtimes of subsequent tests of one condition.

As the number of processes increased, performance improved massively. The speedup was far more than the speedups seen when using MPI.

The work done by each process when the radius was 0 was too trivial to make proper conclusions about scalability from the efficiency table. However, the efficiency table with radius 10 definitely shows the program to be both strongly and weakly scalable. (The efficiency stays constant when the problem size stays constant, and it stays constant when the problem size increases at the same rate as the problem size.) There are outliers in the 32 processor row for 1k, 3k, and 4k, but apart from that, the efficiency stays very close to 1 for the entire table.